

Buffer Overflows



© João Paulo Barraca,
André Zúquete

Security

1

Memory organization topics

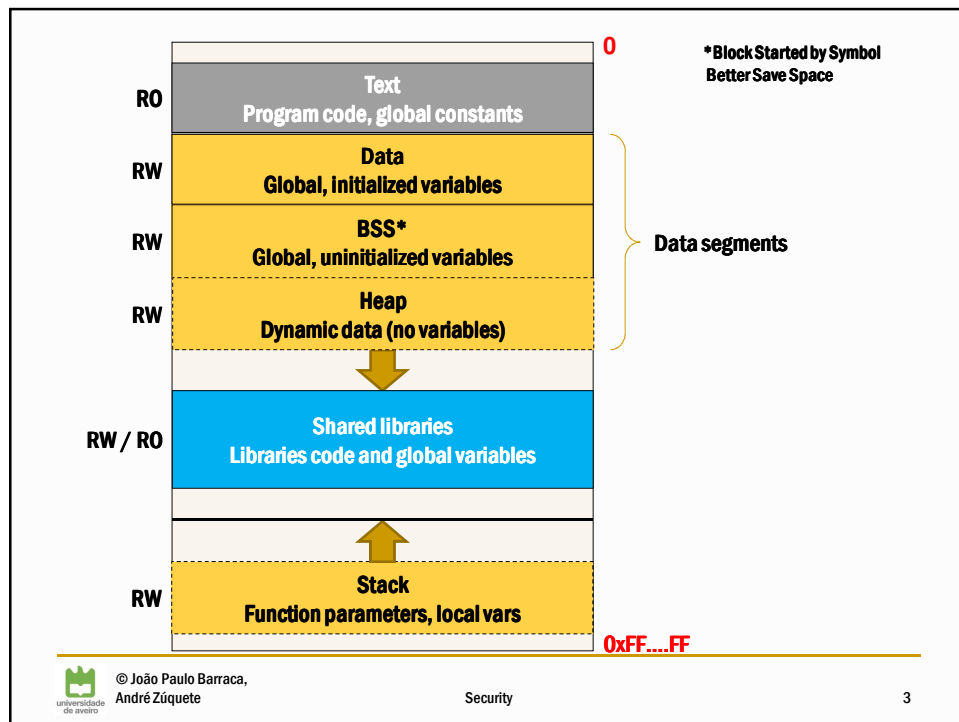
- ▷ Kernel organizes memory in **pages**
 - ♦ Typically 4 kB
- ▷ Processes operate in a **virtual memory space**
 - ♦ Mapped to real 4k pages
 - ♦ Could live in RAM, be file-mapped or be swapped out
- ▷ Kernel groups pages in several segments
 - ♦ **Increases security**
 - Segment-based permissions (RO, RW)
 - ♦ **Increases performance**
 - Some are dynamic: discarded when program terminates
 - Some are static: can be retained, speeding up reuses



© João Paulo Barraca,
André Zúquete

Security

2



mem.c

```
//CONST
const char cntvar[]="constant";

//BSS
static char bssvar[4];

int main(int argc, void** argv)
{
    void * dynmem = malloc(1);
    ...
}
```

```
&main = 0804865c -> text = 08048000
cntvar = 08048920 -> const = 08048000
bssvar = 0804a034 -> bss = 0804a000
&argc = bfeb8590 -> stack = bfeb8000
dynmem = 08435008 -> heap = 08435000
```

mem.c

Content of /proc/self/maps

```
08048000-08049000 r-xp 00000000 08:01 26845750 /home/s/seguranca/mem
08049000-0804a000 r--p 00000000 08:01 26845750 /home/s/seguranca/mem
0804a000-0804b000 rw-p 00001000 08:01 26845750 /home/s/mem
08435000-08456000 rw-p 00000000 00:00 0 [heap]
b7616000-b7617000 rw-p 00000000 00:00 0
b7617000-b776a000 r-xp 00000000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776a000-b776b000 ---p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776b000-b776d000 r--p 00153000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776d000-b776e000 rw-p 00155000 08:01 1574823 /lib/tls/i686/cmov/libc-2.11.1.so
b776e000-b7771000 rw-p 00000000 00:00 0
b777e000-b7782000 rw-p 00000000 00:00 0
b7782000-b7783000 r-xp 00000000 00:00 0 [vdso]
b7783000-b779e000 r-xp 00000000 08:01 1565567 /lib/ld-2.11.1.so
b779e000-b779f000 r--p 0001a000 08:01 1565567 /lib/ld-2.11.1.so
b779f000-b77a0000 rw-p 0001b000 08:01 1565567 /lib/ld-2.11.1.so
bfe99000-bfeba000 rw-p 00000000 00:00 0 [stack]
```



© João Paulo Barraca,
André Zúquete

Security

5

mem.c

Stack evolution:

```
foo [000]: &argc = bfeb8140 -> stack = bfeb8000
foo [001]: &argc = bfdb8110 -> stack = bfdb8000
foo [002]: &argc = bfc8b80e0 -> stack = bfc8b8000
foo [003]: &argc = bfbb80b0 -> stack = bfbb8000
foo [004]: &argc = bfab8080 -> stack = bfab8000
foo [005]: &argc = bf9b8050 -> stack = bf9b8000
foo [006]: &argc = bf8b8020 -> stack = bf8b8000
foo [007]: &argc = bf7b7ff0 -> stack = bf7b7000
foo [008]: &argc = bf6b7fc0 -> stack = bf6b7000
```

Segmentation fault



© João Paulo Barraca,
André Zúquete

Security

6

Some x86 CPU registers

- ▷ General Purpose: A, B, C, D
 - ♦ A: 8bits, AX: 16bits, **EAX: 32bits**, RAX: 64bits
- ▷ BP: Base Pointer (EBP if w/ 32 bits)
 - ♦ Base address of the current function stack frame
 - ♦ A function stack frame is where we have
 - The function parameters
 - The local function variables
- ▷ SP: Stack Pointer (ESP if w/ 32 bits)
 - ♦ Points to end of stack (last value pushed)
- ▷ IP: Instruction Pointer (EIP if w/ 32 bits)
 - ♦ Points to current instruction



© João Paulo Barraca,
André Zúquete


Security

7

Stack segment

```
function ( int arg1, int arg2 )  
{  
    int var1 = arg1;  
    int var2;  
}
```

- ▷ Stack is used to
 - ♦ Pass parameters to functions (eg. **arg1**)
 - ♦ Store local variables (eg. **var1**)
- ▷ Values are PUSHed or POPed from stack
 - ♦ eg: **push eax, pop eax**
- ▷ Allocation of local variables in space
 - ♦ **int var1; → sub esp, 4**
- ▷ Accessing variables in the stack
 - ♦ A parameter:
 - $\text{arg1} \rightarrow \text{ebp} + 8$
 - $\text{arg2} \rightarrow \text{ebp} + 12$
 - ♦ A local variable:
 - $\text{var1} \rightarrow \text{ebp} - 4$
 - $\text{var2} \rightarrow \text{ebp} - 8$




© João Paulo Barraca,
André Zúquete

Security

8

Initialization of a stack frame

- ▷ This is done in the **prologue** of a function
 - ♦ And is undone at its **epilogue**
- ▷ The prologue consists in:
 - ♦ Saving the base address of the previous stack frame and setting the new one
 - `push ebp`
 - `mov ebp, esp`
 - ♦ Allocate space for local variables
 - `sub esp, Imm`
- ▷ The epilogue is
 - `mov esp, ebp`
 - `pop ebp` } **leave**




© João Paulo Barraca,
André Zúquete

Security

9

Function call and return

- ▷ Call steps
 - ♦ Put arguments in stack
 - Usually with `PUSH`
 - ♦ Call the function address
 - Pushes the IP to the stack (**return address**)
 - IP has the **next instruction address**
 - ♦ Release stack space
 - Usually increasing `ESP`
- ▷ Returning from a function
 - ♦ The `RET` instruction pops the saved IP (return address)



© João Paulo Barraca,
André Zúquete


Security

10

Function foo()

```
void foo ( int arg1, int arg2 )    foo:
{
    int var1 = arg1;
    int var2;
}

    push ebp
    mov ebp, esp
    sub esp, 8
    mov eax, DWORD PTR [ebp+8]
    mov DWORD PTR [ebp-4], eax
    leave
    ret
```



© João Paulo Barraca,
André Zúquete

Security

11

Buffer overflow

- ▷ Write beyond the boundaries of a buffer
- ▷ Consequences
 - ♦ Write over other values located next to the buffer
 - ♦ Write over special values co-located (saved registers)
 - Saved BP
 - Damages the base address of the previous stack frame
 - Saved IP (return address)
 - Jump to any address on return!



© João Paulo Barraca,
André Zúquete

Security

12


Stack smashing attack

▷ Roadmap

- ♦ Overflow a local variable
- ♦ Extend the overflow to the return address
- ♦ Change the return address in order to jump to the injected data
 - Which should be executable code
- ♦ Wait for the return of the function

▷ Difficulty

- ♦ A return using a saved address is an absolute jump
- ♦ The attacker needs to know the absolute address of the vulnerable variable
 - Given the source code, knowing the machine and the initial stack address, this is feasible



© João Paulo Barraca,
André Zúquete

Security

13

bo.c

```
int foo()
{
    char a[4];
    scanf("%s", a);
}
```

Pre-allocation of space for function call parameters in advance (and excess)

Allows function calls without pushing/poping values to/from the stack

```
.LC0:
.string "%s"
.text

foo:
    push ebp
    mov ebp, esp
    sub esp, 40
    mov eax, OFFSET FLAT:.LC0
    lea edx, [ebp-12]
    mov DWORD PTR [esp+4], edx
    mov DWORD PTR [esp], eax
    call __isoc99_scanf
    leave
    ret
```



© João Paulo Barraca,
André Zúquete

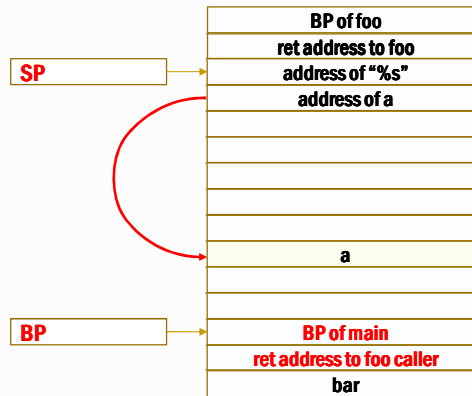
Security

14

bo.s

```
.LC0:
.string "%s"
.text

foo:
push ebp
mov ebp, esp
sub esp, 40
mov eax, OFFSET FLAT:.LC0
lea edx, [ebp-12]
mov DWORD PTR [esp+4], edx
mov DWORD PTR [esp], eax
call __isoc99_scanf
leave
ret
```



© João Paulo Barraca,
André Zúquete

Security

15

Buffer overflow

```
[jpbarraca@atnog: seguranca]$ ./bo a Write inside a
[jpbarraca@atnog: seguranca]$ ./bo aa Write inside a
[jpbarraca@atnog: seguranca]$ ./bo aaaaaaaaaa Write outside a
[jpbarraca@atnog: seguranca]$ ./bo aaaaaaaaaa Write over stored BP
Segmentation fault
```



© João Paulo Barraca,
André Zúquete

Security

16

Mitigation:


Prevention mechanisms

- ▷ Avoid execution of injected instructions
 - ♦ In segments/pages that usually have no code
 - ♦ Prevents the execution of code injected as data
- ▷ Randomize the address space
 - ♦ ADLR (Address Space Layout Randomization)
 - ♦ Segments do not start in fixed positions on each run of the same application
 - But segments keep their relative position
 - ♦ Prevents jumps to well-known code locations



Mitigation:

Prevention mechanisms



- ▷ Variable reordering
 - ♦ Usually the vulnerable variables are arrays
 - ♦ To protect other kinds of local variables (in the same stack frame), arrays are moved closer to the saved registers
 - ♦ This reduces the set of variables that may be affected by a buffer overrun



Mitigation: Detection mechanisms

▷ Stack canaries

- ♦ A value unknown to attackers (canary) is stored next to saved registers
 - Saved BP and return address
- ♦ Stack smashing attacks usually cannot affect saved registers with running over a canary
 - Because they are usually based on string overruns
- ♦ The canary is checked before the function's epilogue
 - If different from the original value, an exception is raised

