

Database security

Advantages of using databases

- ♦ Shared access
 - Many users use one common, centralized data set
- ♦ Minimal redundancy
 - Individual users do not have to collect and maintain their own sets of data
- ♦ Data consistency
 - A change to a data value affects all users of the data value
- ♦ Data integrity
 - Data values are protected against accidental or malicious undesirable changes
- ♦ Controlled access
 - Only authorized users are allowed to view or to modify data values

Security requirements (1/2)

- ♦ Physical integrity
 - Immunity to physical problems
 - e.g. power failures
 - Ability to reconstruct the database if destroyed in a catastrophe
- ♦ Logical integrity
 - Structure is preserved
- ♦ Element integrity
 - Data in each element is accurate
- ♦ Auditability
 - It is possible to track who or what has accessed (or modified) which elements in the database

Security requirements (2/2)

- ♦ Access control
 - A user/role is allowed to access only authorized data/queries
 - Different users/roles can be restricted to different modes
 - e.g. read or write
- ♦ User authentication
 - Every user/role is positively identified
 - Fundamental for audit trails and for permissions to access data
- ♦ Availability
 - Users/roles can access the database in general and all the data for which they are authorized

Two-phase updates

- ♦ Problem

- Failures during updates may render databases incoherent
 - Requirement → atomicity (from ACID)
- Logical integrity problem

- ♦ Solution: Atomicity with two-phase updates

Two-phase update

- ♦ **1st phase: intent phase**
 - The DBMS gathers resources it needs to perform the update
 - It does everything to prepare for the update, but makes no changes to the database
 - Committing: writes a commit flag to the database
 - Point of no return
 - After, the DBMS begins making permanent changes
- ♦ **2nd phase: commit phase**
 - Makes the permanent changes in the database
 - It lasts until finishing all changes prepared in the first phase
 - When it finishes, the database changed to a new stable, coherent state

Redundancy / internal consistency

- ♦ Error detection and correction codes
 - Parity bits, Hamming codes, cyclic redundancy checks
 - Can be applied to different data elements
 - Fields, records, entire database
 - More space
 - To store error detection/correction information
- ♦ Shadow fields
 - Duplication of fields or records
 - Requires substantial storage space

Concurrency / consistency

- ♦ Accesses by two users of the same DBMS must be constrained so that neither interferes with each other
 - Simple locking: multiple readers, one writer
 - But simple locking may not be enough on query-update cycles
- ♦ Solution
 - Treat every query-update cycle as a single atomic operation (a transaction)
 - e.g. flight booking
 - Synchronization should be applied to transactions
 - Two concurrent transactions cannot write (and sometimes read) the same field/record

Monitors

- ♦ DBMS unit responsible for the DB structural integrity
 - Checks entered values to ensure their consistency with field, record or database consistency constraints
- ♦ Types of monitors
 - Range comparisons
 - Tests if values belong to an acceptable range
 - State constraints
 - Describe the condition of the entire database
 - e.g. the commit flag
 - Impose integrity restriction rules
 - e.g. to detect duplicate records
 - Transition constraints
 - Describe required conditions before changing the database

Database activity monitoring

- ♦ DBMS usage supervision
 - To detect abuses
 - To detect unusual/suspect activity or operations
- ♦ DBMS independent
 - Not part of the DBMS
 - External observation of DBMS activity
- ♦ Monitoring sensors
 - Network activity
 - Local SQL commands performed
 - Log analysis

Sensitive data

- ♦ Data requiring (extra) protection
 - From loss (disclosure)
 - From misuse
 - From modification
 - From unauthorized access
- ♦ Risks
 - Privacy and welfare of individuals
 - Business activities
 - Security-related activities

Sensitive data

- ♦ Some databases contain sensitive data
 - Data that should not be made public
 - e.g. clinical records of patients
- ♦ Sensitivity depends on: BD purpose + DB data
 - Some record fields, entire records/tables, entire database
 - e.g. personal health record (HER) with all detected pathologies, treatments and interventions
 - e.g. clinical records of an AIDS table
 - e.g. defense-related databases
- ♦ Complexity
 - Simple cases: all or nothing
 - Everything sensitive, nothing sensitive
 - Complicated cases: part of the DB elements are sensitive
 - In some cases, sensitivity is extended to the simple existence of a field data or record

Sensitive data:

Factors that make data sensitive

- ♦ Inherently sensitive
 - The value itself may be so revealing that it is sensitive
- ♦ From a sensitive source
 - The value may reveal the identity of its source
- ♦ Declared sensitive
 - The value was explicitly declared sensitive
- ♦ Belongs to a sensitive record
 - Value of a record was explicitly marked as sensitive
- ♦ Sensitive given previously disclosed information
 - By itself, the data is not sensitive, but together with other data, the whole can be sensitive

Types of disclosures (of sensitive data)

- ♦ Exact data
 - The exact value of a sensitive datum
 - The most serious disclosure
- ♦ Bounds
 - Sensitive data item is $>$ lower bound or $<$ upper bound
 - Sometimes bounds are used to protect (hide) sensitive data
 - By providing bounds to elements instead of their exact value
- ♦ Negative result
 - By getting a negative result for a query on a sensitive value, a user can conclude that the value has a particular set of values
 - e.g. from a list of effective voters we can conclude who didn't vote

Types of disclosures (of sensitive data)

- ♦ Existence

- The existence of a sensitive field in a record can be, by itself, sensitive information
 - Because it may reveal a hidden data gathering & processing activity

- ♦ Probable value

- By crossing the results of several queries we can infer a probability for an element value

Inference

- ♦ Definition
 - A way to infer, or derive, sensitive data from non-sensitive data

Inference attacks

- ♦ Direct attack

- Uses queries with a blend of selection rules that use sensitive fields and non-sensitive fields
- The DBMS can be deceived by the selection rules with non-sensitive fields, which are not intended to select particular records

- ♦ Indirect attack

- Inference of particular values from statistical values computed over several records
 - Counts, sums, averages

Inference attacks

- ♦ Tracker attack

- The database may conceal data when a small number of records make up the large proportion of the data revealed
- A tracker attack can fool the DBMS by using different queries that reveal data and, by combining the results, the attacker can get the desired information

K-anonymity

L. Sweeney, “K-anonymity: A Model for Protecting Privacy”, Int. Journal on Uncertainty, Fuzziness and Knowledge-based Systems. 2002

- ♦ Definition

- No query can deliver an **anonymity set** with less than **k** entries
- The **anonymity set** is the set of all possible subjects

- ♦ Privacy-critical attributes

- (Unique) identifiers
- Quasi-identifiers
 - When combined can produce unique tuples
- Sensitive attributes
 - Potentially unique per subject
 - Disease, salary, crime committed

Multilevel security:

Goal

- ♦ Tag information items with security classifications
 - e.g unclassified, confidential, secret, top secret
- ♦ Tag queries with security levels
 - The security level of the entity responsible for the query
- ♦ Prevent queries from observing values of fields with a different security classification
 - Or from observing meaningful values

Multilevel security:

Poli-instanciación

- ♦ A record with a particular key field may be duplicated in different security levels
 - Possibly with different values
- ♦ This reduces the precision of the database information
 - The correctness of the information depends on the entity performing the query
 - Duplicates know can legitimately occur

Multilevel security:

Separation strategies (1)

- ♦ Partitioning
 - Different security levels, different databases
 - Queries are directed to the appropriate DB
- ♦ Advantages
 - Easy to implement
- ♦ Disadvantages
 - Redundancy of information
 - Problems in the access to records with fields with different security levels

Multilevel security:

Separation strategies (2)

- ♦ Encryption
 - Fields are encrypted with a security-level key
- ♦ Advantages
 - Single database, same database structure
- ♦ Disadvantages
 - Decryption on each query with the adequate security level key
 - Randomized encryption: equal fields should not produce the same cryptogram
 - Otherwise statistics and known-plaintext attacks disclose values
 - Solution: different keys per record or different IVs per record
 - No encrypted values should be updated by providing another encrypted value

Multilevel security:

Separation strategies (3)

- ♦ Integrity lock
 - Each data item is formed by three parts:
 - Data item, sensitivity label, checksum
 - The sensitivity label should be
 - Unforgeable (cannot be changed)
 - Unique (cannot be copied to another data item)
 - Concealed (cannot be observed)
- ♦ Advantages
 - Can use a regular DBMS
 - Trusted stored procedures are enough to implement them
- ♦ Disadvantages
 - Space for storing sensitivity labels and checksums

Laws for the protection of personal data

- ♦ Each country has its own set of laws
 - There is not a global consensus
- ♦ In Portugal this is supervised by CNPD
 - **Comissão Nacional de Proteção de Dados**
 - All data processing involving personal data gathered from individuals needs to be submitted to CNPD for authorization
- ♦ European Directive for Data Protection
 - To be applied from May 25, 2018