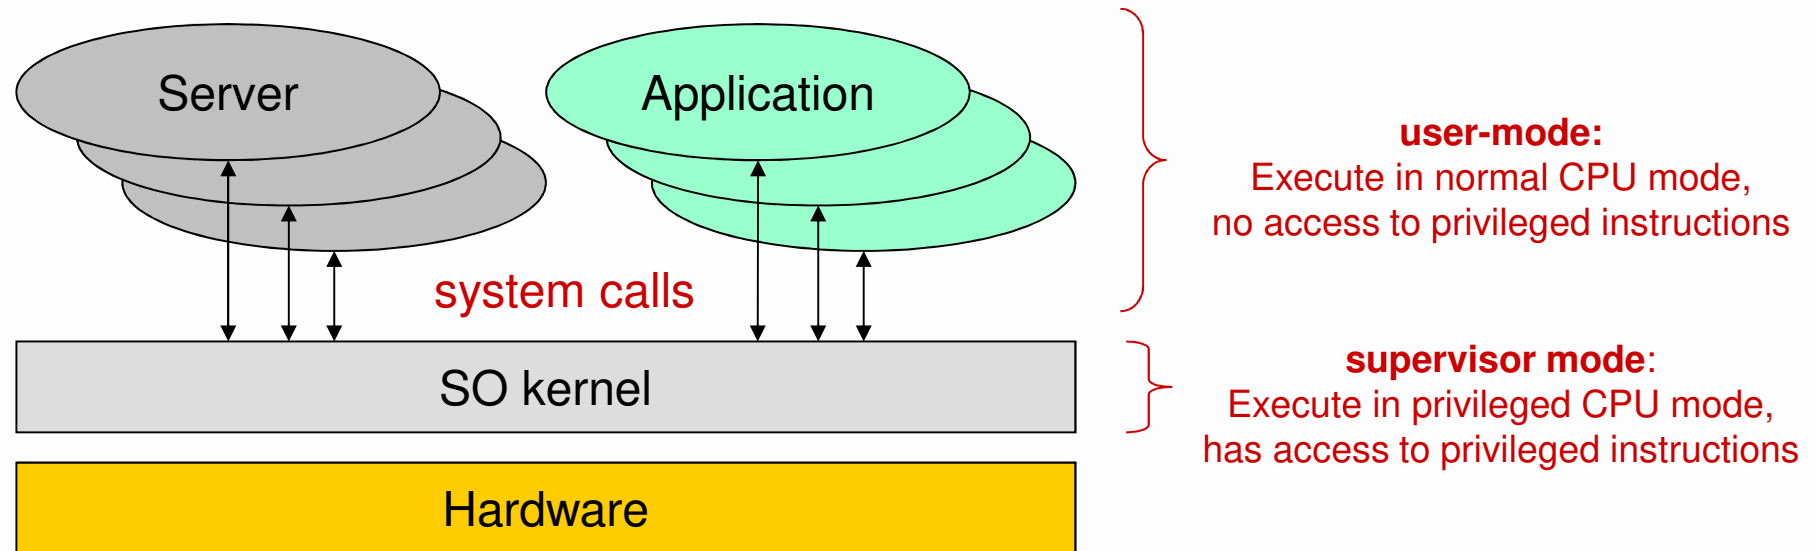

Security in Operating Systems


Operating system



- ♦ Kernel mission

- Virtualize the hardware
 - Computational model
- Enforce protection policies and provide protection mechanisms
 - Against involuntary mistakes
 - Against non-authorized activities

Execution rings



- ♦ Different levels of privilege
 - Forming a set of concentric rings
 - Used by CPU's to prevent non-privileged code from running privileged opcodes
 - e.g. IN/OUT, TLB manipulation
- ♦ Nowadays processors have 4 rings
 - But OS's usually use only 2
 - 0 (supervisor/kernel mode) and 3 (user-mode)
- ♦ Transfer of control between rings requires special gates
 - The ones that are used by syscalls

Execution of virtual machines


- ♦ Common approach

- Software-based virtualization
- Direct execution of guest user-mode code
- Binary translation of privileged code
 - Guest OS kernels remain unchanged, but do not run directly on the host machine

- ♦ Hardware-assisted virtualization

- Full virtualization
- There is a ring -1 below ring 0
 - Hypervisor
- It can virtualize hardware for many ring 0 kernels
 - No need of binary translation
 - Guest OS's run faster

Computational model



- ♦ Set of entities (objects) managed by the OS kernel
 - User identifiers
 - Processes
 - Virtual memory
 - Files and file systems
 - Communication channels
 - Physical devices
 - Storage
 - Magnetic disks, optical disks, silicon disks, tapes
 - Network interfaces
 - Wired, wireless
 - Human-computer interfaces
 - Keyboards, graphical screens, text consoles, mice
 - Serial/parallel I/O interfaces
 - USB, serial ports, parallel ports, infrared, bluetooth

Computational model:

User identifiers



- For the OS kernel a user is a number
 - Established during a login operation
 - User ID (UID)
- All activities are executed on a computer on behalf of a UID
 - The UID allows the kernel to assert what is allowed/denied to processes
 - Linux: UID 0 is omnipotent (root)
 - Administration activities are usually executed with UID 0
 - Windows: concept of privileges
 - For administration, system configuration, etc.
 - There is no unique, well-known identifier for and administrator
 - Administration privileges can be bound to several UIDs
 - Usually through administration groups
 - Administrators, Power Users, Backup Operators

Computational model:

Group identifiers



- ♦ Groups also have an identifier
 - A group is a set of users
 - A group can be defined by including other groups
 - Group ID (GID)
- ♦ A user can belong to several groups
 - Rights = UID rights + rights of his groups
- ♦ In Linux all activities are executed on behalf of a set of groups
 - Primary group
 - Typically used for setting file protection
 - Secondary groups

Computational model:

Processes

- ♦ A process defines the context of an activity
 - For taking security-related decisions
 - For other purposes (e.g. scheduling)
- ♦ Security-related context
 - Identity (UID and GIDs)
 - Fundamental for enforcing access control
 - Resources being used
 - Open files
 - Including communication channels
 - Reserved virtual memory areas
 - CPU time used

Access control

- ♦ The OS kernel is an access control monitor
 - Controls all interactions with the hardware
 - Controls all interactions between entities of the computational model
- ♦ Subjects
 - Usually local processes
 - Through the syscall API
 - A syscall is not an ordinary function call
 - But also messages from other hosts

Mandatory access controls

- ♦ OS kernels are loaded of mandatory access control policies
 - They are part of the computational model logic
 - They cannot be overruled not even by administrators
 - Unless they change the OS kernel behavior
- ♦ Examples:
 - Kernel runs in CPU privileged modes, user applications run in non-privileged modes
 - Separation of virtual memory areas
 - Inter-process signaling
 - Interpretation of files' ACLs

Protection with ACLs

- ♦ Each object has an ACL
 - It says which subjects can do what
- ♦ An ACL can be discretionary or mandatory
 - When mandatory it cannot be modified
 - When discretionary it can be tailored
- ♦ An ACL is checked when an activity, on behalf of a subject, wants to manipulate the object
 - Ifs the manipulation request is not authorized by the ACL, the access is denied
 - The SO kernel is the responsible for enforcing ACL-based protection
 - It acts as a security monitor

Protection with capabilities

- ♦ Less common in normal OS kernels
 - Though there are some good examples
- ♦ Example: open file descriptors
 - Applications' processes indirectly manipulate file descriptors through the OS kernel
 - The OS kernel has full control over file descriptors
 - File descriptors can only be granted to other processes through the OS kernel
 - Changes in the protection of files does not impact existing file descriptors

Unix file protection ACLs:

Fixed-structure, discretionary ACL

- ♦ Each file system object has an ACL
 - Binding 3 rights to 3 subjects
 - Only the owner can update the ACL
- ♦ Rights: **R W X**
 - Read right / Listing right
 - Write right / create or remove files or subdirectories
 - Execution right / use as process' current working directory
- ♦ Subjects:
 - An UID (owner)
 - A GID
 - Others

Windows NTFS file protection:

Variable-size, discretionary ACLs

- ♦ Each file system object has an ACL and a owner
 - The ACL grants 14 types of access rights to a variable-size list of subjects
 - Owner can be an UID or a GID
 - Owner has no special rights over the ACL
- ♦ Subjects:
 - Users (UIDs)
 - Groups (GIDs)
 - The group "Everyone" stands for anybody
- ♦ Rights:
 - Traverse Folder / Execute File
 - List Folder / Read Data
 - Read Attributes
 - Read Extended Attributes
 - Create Files /Write Data
 - Create Folders / Append Data
 - Write Attributes
 - Write Extended Attributes
 - Delete Subfolders and Files
 - Delete
 - Read Permissions
 - Change Permissions
 - Take Ownership

Privilege elevation:

Set-UID mechanism

- ♦ It is used to change the UID of a process running a program stored on a Set-UID file
 - If the program file is owned by UID X and the set-UID ACL bit is set, then it will be executed in a process with UID X, independently of the UID of the subject that executed the program
- ♦ It is used to provide privileged programs for running administration task invoked by normal, untrusted users
 - Change the user's password (passwd)
 - Change to super-user mode (su, sudo)
 - Mount devices (mount)

Privilege elevation:

Set-UID mechanism (cont.)

- ♦ Effective UID / Real UID
 - Real UID is the UID of the process creator
 - App launcher
 - Effective UID is the UID of the process
 - The one that really matters for defining the rights of the process
- ♦ UID change
 - Ordinary application
 - eUID = rUID = UID of process that executed **exec**
 - eUID cannot be changed (unless = 0)
 - Set-UID application
 - eUID = UID of **exec**'d application file, rUID = initial process UID
 - eUID can revert to rUID
 - rUID cannot change

Privilege elevation:

sudo mechanism

- ♦ Administration by root is not advised
 - One "identity", many people
 - Who did what?
- ♦ Preferable approach
 - Administration role (uid = 0), many users assume it
 - Sudoers
 - Defined by a configuration file used by sudo
- ♦ sudo is a Set-UID application with UID = 0
 - Appropriate logging can take place on each command run with sudo

Privilege reduction: chroot mechanism (or jail)

- ♦ Used to reduce the visibility of a file system
 - Each process descriptor has a root i-node number
 - From which absolute pathname resolution takes place
 - chroot changes it to an arbitrary directory
 - The process' file system view gets reduced
- ♦ Used to protect the file system from potentially problematic applications
 - e.g. public servers, downloaded applications
 - But it is not bullet proof!

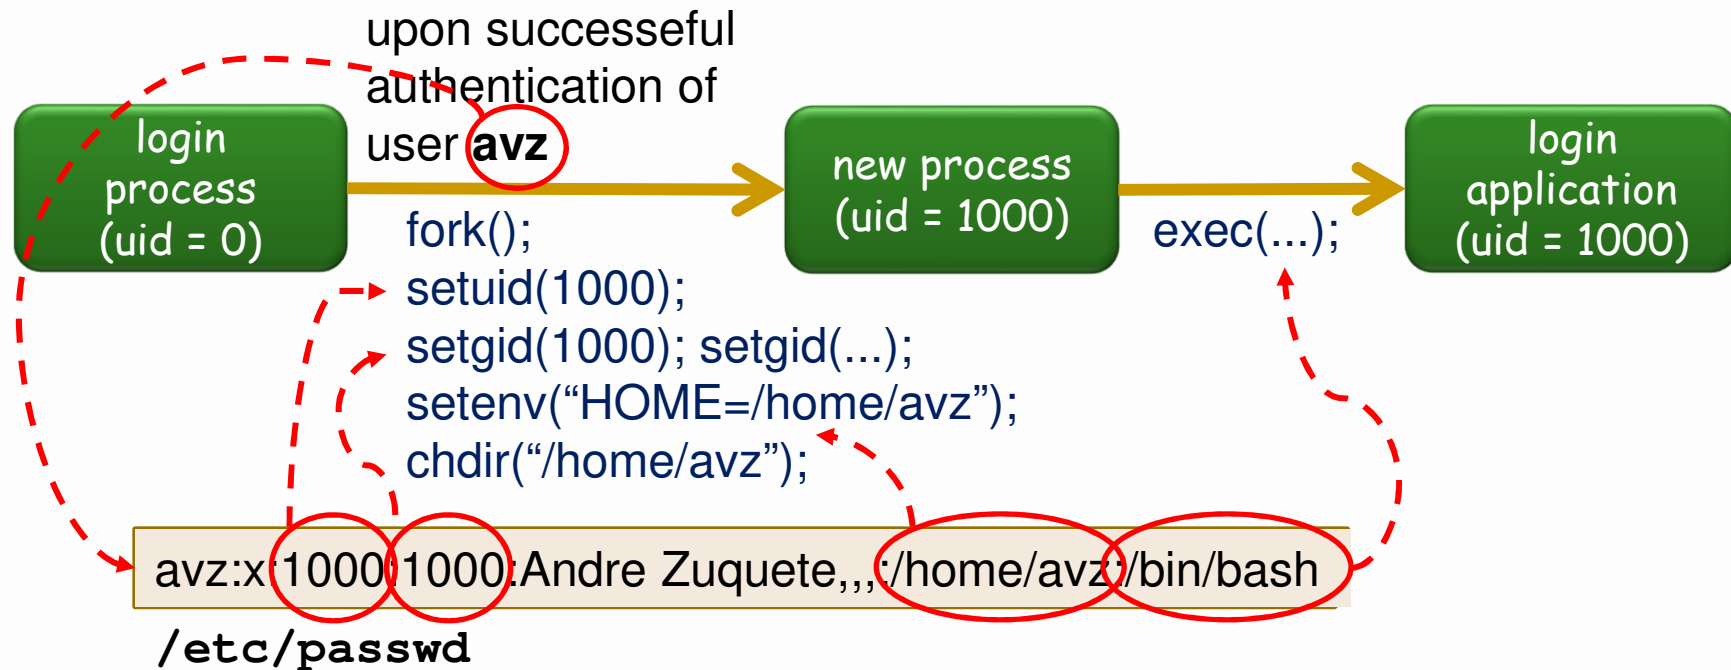
Linux login:

Not an OS kernel operation

- ♦ A privileged login application presents a login interface for getting users' credentials
 - A username/password pair
 - Biometric data
 - Smartcard and activation PIN
- ♦ The login application validates the credentials and fetches the appropriate UID and GIDs for the user
 - And starts an initial user application on a process with those identifiers
 - In a Linux console this application is a shell
 - When this process ends the login application reappears
- ♦ Thereafter all processes created by the user have its identifiers
 - Inherited through forks

Linux: from login to session processes

- ♦ The login process must be a privileged process
 - Has to create processes with arbitrary UID and GIDs
 - The ones of the entity logging in



Login in Linux:

Password validation process

- Username is used to fetch a UID/GID pair from `/etc/passwd`
 - And a set of additional GIDs in the `/etc/group` file
- Supplied password is transformed using a digest function
 - Currently configurable, for creating a new user (`/etc/login.conf`)
 - Its identification is stored along with the transformed password
- The result is checked against a value stored in `/etc/shadow`
 - Indexed again by the username
 - If they match, the user was correctly authenticated
- File protections
 - `/etc/passwd` and `/etc/group` can be read by anyone
 - `/etc/shadow` can only be read by root
 - Protection against dictionary attacks